

Nikshay Jain - MM21B044

```
import math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.experimental import enable_hist_gradient_boosting
from sklearn.ensemble import HistGradientBoostingClassifier

from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

c:\Users\Nikshay Jain\AppData\Local\Programs\Python\Python38\lib\site-
packages\sklearn\experimental\enable_hist_gradient_boosting.py:16:
UserWarning: Since version 1.0, it is not needed to import
enable_hist_gradient_boosting anymore. HistGradientBoostingClassifier
and HistGradientBoostingRegressor are now stable and can be normally
imported from sklearn.ensemble.
  warnings.warn(

columns = ["parents", "has_nurs", "form", "children", "housing",
"finance", "social", "health", "class"]
data = pd.read_csv('nursery.data', names=columns)
data
```

	parents	has_nurs	form	children	housing
finance \					
0	usual	proper	complete	1	convenient
convenient					
1	usual	proper	complete	1	convenient
convenient					
2	usual	proper	complete	1	convenient
convenient					
3	usual	proper	complete	1	convenient
convenient					
4	usual	proper	complete	1	convenient
convenient					
...
.					..
12955	great_pret	very_crit	foster	more	critical

```

inconv
12956 great_pret very_crit foster more critical
inconv
12957 great_pret very_crit foster more critical
inconv
12958 great_pret very_crit foster more critical
inconv
12959 great_pret very_crit foster more critical
inconv

```

```

          social      health      class
0      nonprob  recommended  recommend
1      nonprob    priority    priority
2      nonprob  not_recom  not_recom
3  slightly_prob  recommended  recommend
4  slightly_prob    priority    priority
...
12955  slightly_prob    priority  spec_prior
12956  slightly_prob  not_recom  not_recom
12957  problematic  recommended  spec_prior
12958  problematic    priority  spec_prior
12959  problematic  not_recom  not_recom

```

```
[12960 rows x 9 columns]
```

```

# Encode the target variable
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(data['class'])

```

Decision tree - categorical features

```

# Convert categorical columns to indices (integer-based for
HistGradientBoosting)
X = data.drop(columns=["class"])

# Convert categorical columns to indices
for col in X.columns:
    X[col] = pd.Categorical(X[col]).codes

# Split data into train, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(X, y,
test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp,
test_size=0.75, random_state=42)

hgb_params = {
    'max_iter': [100, 300],
    'learning_rate': [0.1, 0.2, 0.5],
    'max_depth': [2, 3, 5],
    'min_samples_leaf': [1, 2],

```

```

}

# HistGradientBoostingClassifier with categorical support
categorical_features = np.arange(X_train.shape[1])
clf_hgb =
HistGradientBoostingClassifier(categorical_features=categorical_features,
random_state=42)
clf_hgb = GridSearchCV(clf_hgb, hgb_params, cv=5, scoring='accuracy',
n_jobs=-1, verbose=2)
clf_hgb.fit(X_train, y_train)

# Get the best parameters
print(f"Best Parameters for categorical features in Decision tree:
{clf_hgb.best_params_}")

# Evaluate the model
y_val_pred_hgb = clf_hgb.predict(X_val)
hgb_acc = accuracy_score(y_val, y_val_pred_hgb)
print(f"Decision Tree (Categorical Features) Val Accuracy: {hgb_acc}")

# Evaluate on test set
y_test_pred_hgb = clf_hgb.predict(X_test)
test_acc_hgb = accuracy_score(y_test, y_test_pred_hgb)
print(f"Decision Tree (Categorical Features) Test Accuracy:
{test_acc_hgb}")

Fitting 5 folds for each of 36 candidates, totalling 180 fits
Best Parameters for categorical features in Decision tree:
{'learning_rate': 0.1, 'max_depth': 5, 'max_iter': 300,
'min_samples_leaf': 1}
Decision Tree (Categorical Features) Val Accuracy: 1.0
Decision Tree (Categorical Features) Test Accuracy: 0.9993141289437586

```

Decision tree - one hot encoding

```

X = data.drop(columns=["class"])

# One hot encoding for features
ohe = OneHotEncoder(sparse=False)
X = ohe.fit_transform(X)

# Split data into train, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(X, y,
test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp,
test_size=0.75, random_state=42)

c:\Users\Nikshay Jain\AppData\Local\Programs\Python\Python38\lib\site-
packages\sklearn\preprocessing\_encoders.py:808: FutureWarning:
`sparse` was renamed to `sparse_output` in version 1.2 and will be

```

removed in 1.4. ``sparse_output`` is ignored unless you leave ``sparse`` to its default value.

```
warnings.warn(
```

```
dt_params = {
    'max_depth': [5, 10, 15, 20],
    'min_samples_split': [1, 2, 5],
    'min_samples_leaf': [1, 2, 4]
}

# Decision Tree on One-Hot Encoded Data
clf_dt_onehot = GridSearchCV(DecisionTreeClassifier(random_state=42),
dt_params, cv=5)
clf_dt_onehot.fit(X_train, y_train)
print(f"Best Decision Tree (One-Hot) Params:
{clf_dt_onehot.best_params_}")

y_val_pred_dt = clf_dt_onehot.predict(X_val)
dt_val_acc = accuracy_score(y_val, y_val_pred_dt)
print(f"Decision Tree (One-Hot Encoded) Val Accuracy: {dt_val_acc}")

y_test_pred_dt = clf_dt_onehot.predict(X_test)
dt_test_acc = accuracy_score(y_test, y_test_pred_dt)
print(f"Decision Tree (One-Hot Encoded) Test Accuracy: {dt_test_acc}")

Best Decision Tree (One-Hot) Params: {'max_depth': 15,
'min_samples_leaf': 1, 'min_samples_split': 1}
Decision Tree (One-Hot Encoded) Val Accuracy: 0.9969135802469136
Decision Tree (One-Hot Encoded) Test Accuracy: 0.9948559670781894
```

Logistic Regression

```
logreg_params = {
    'C': [1, 10, 100],
    'solver': ['liblinear'],
    'max_iter': [50, 100, 200]
}

clf_lr = GridSearchCV(LogisticRegression(penalty='l1',
random_state=42), logreg_params, cv=5)
clf_lr.fit(X_train, y_train)

y_val_pred_lr = clf_lr.predict(X_val)
lr_val_acc = accuracy_score(y_val, y_val_pred_lr)

print(f"Best Logistic Regression Params: {clf_lr.best_params_}")
print(f"Logistic Regr (L1 Regularization) Val Accuracy: {lr_val_acc}")

y_test_pred_lr = clf_lr.predict(X_test)
lr_test_acc = accuracy_score(y_test, y_test_pred_lr)
```

```
print(f"Logistic Regr (L1 Regularization) Test Accuracy: {lr_test_acc}")
```

Best Logistic Regression Params: {'C': 100, 'max_iter': 50, 'solver': 'liblinear'}

Logistic Regr (L1 Regularization) Val Accuracy: 0.9022633744855967

Logistic Regr (L1 Regularization) Test Accuracy: 0.9245541838134431

K Nearest Neighbours

```
# k-Nearest Neighbors (kNN)
```

```
clf_knn = KNeighborsClassifier()
```

```
param_grid_knn = {'n_neighbors': [3, 7, 11, 15],  
                  'weights': ['uniform', 'distance']}
```

```
}
```

```
grid_search_knn = GridSearchCV(clf_knn, param_grid_knn, cv=5)
```

```
grid_search_knn.fit(X_train, y_train)
```

```
print(f"Best kNN Params: {grid_search_knn.best_params_}")
```

```
y_val_pred_knn = grid_search_knn.predict(X_val)
```

```
knn_val_acc = accuracy_score(y_val, y_val_pred_knn)
```

```
print(f"k-Nearest Neighbors Val Accuracy: {knn_val_acc}")
```

```
y_test_pred_knn = grid_search_knn.predict(X_test)
```

```
knn_test_acc = accuracy_score(y_test, y_test_pred_knn)
```

```
print(f"k-Nearest Neighbors Test Accuracy: {knn_test_acc}")
```

Best kNN Params: {'n_neighbors': 11, 'weights': 'distance'}

k-Nearest Neighbors Val Accuracy: 0.9526748971193416

k-Nearest Neighbors Test Accuracy: 0.9639917695473251

```
n_repeats = 5
```

```
methods = ['Decision Tree (Categorical)', 'Decision Tree (One-Hot)',  
           'Logistic Regression (L1)', 'k-NN']
```

```
accuracies = {method: [] for method in methods}
```

```
for _ in range(n_repeats):
```

```
    X = data.drop(columns=["class"])
```

```
    # Convert categorical columns to indices
```

```
    for col in X.columns:
```

```
        X[col] = pd.Categorical(X[col]).codes
```

```
    X_train_temp, X_val_temp, y_train_temp, y_val_temp =  
    train_test_split(X, y, test_size=0.1)
```

```
    y_val_pred_dt_cat = clf_hgb.predict(X_val_temp)
```

```
    accuracies['Decision Tree  
(Categorical)'].append(accuracy_score(y_val_temp, y_val_pred_dt_cat))
```

```

# One hot encoded
X = data.drop(columns=["class"])
X = ohe.fit_transform(X)

X_train_temp1, X_val_temp1, y_train_temp1, y_val_temp1 =
train_test_split(X, y, test_size=0.1)

y_val_pred_dt_oh = clf_dt_onehot.predict(X_val_temp1)
accuracies['Decision Tree (One-
Hot)'].append(accuracy_score(y_val_temp1, y_val_pred_dt_oh))

y_val_pred_lr = clf_lr.predict(X_val_temp1)
accuracies['Logistic Regression
(L1)'].append(accuracy_score(y_val_temp1, y_val_pred_lr))

y_val_pred_knn = grid_search_knn.predict(X_val_temp1)
accuracies['k-NN'].append(accuracy_score(y_val_temp1,
y_val_pred_knn))

# Calculate mean and variance of accuracies
mean_accuracies = {method: np.mean(accuracies[method]) for method in
methods}
var_accuracies = {method: np.var(accuracies[method]) for method in
methods}

# Visualization
x = np.arange(len(methods))
means = [mean_accuracies[method] for method in methods]
variances = [var_accuracies[method] for method in methods]

# extracted from the website of the database
additional_methods = ['XGBoost', 'SVC', 'Random Forest', 'Neural
Network', 'Logistic Regression (L2)']
additional_means = [0.99969, 0.99198, 0.98025, 1.0, 0.92253]
additional_vars = [0.00001, 0.00385, 0.00453, 0, 0.00895]

methods.extend(additional_methods)
means = [mean_accuracies[method] for method in mean_accuracies] +
additional_means
variances = [var_accuracies[method] for method in var_accuracies] +
additional_vars

plt.figure(figsize=(10, 6))
plt.errorbar(means, methods, xerr=variances, fmt='o', capsize=5,
capthick=2, marker='s', markersize=8, label='Mean Accuracy')

# Adding labels and title
plt.ylabel('Models')
plt.xlabel('Accuracy')

```

```
plt.title('Model Accuracy with Error Bars (Mean & Variance)')
plt.grid(True)
```

```
# Show the plot
```

```
plt.tight_layout()
plt.show()
```

```
c:\Users\Nikshay Jain\AppData\Local\Programs\Python\Python38\lib\site-
packages\sklearn\preprocessing\_encoders.py:808: FutureWarning:
`sparse` was renamed to `sparse_output` in version 1.2 and will be
removed in 1.4. `sparse_output` is ignored unless you leave `sparse`
to its default value.
```

```
warnings.warn(
```

```
c:\Users\Nikshay Jain\AppData\Local\Programs\Python\Python38\lib\site-
packages\sklearn\preprocessing\_encoders.py:808: FutureWarning:
`sparse` was renamed to `sparse_output` in version 1.2 and will be
removed in 1.4. `sparse_output` is ignored unless you leave `sparse`
to its default value.
```

```
warnings.warn(
```

```
c:\Users\Nikshay Jain\AppData\Local\Programs\Python\Python38\lib\site-
packages\sklearn\preprocessing\_encoders.py:808: FutureWarning:
`sparse` was renamed to `sparse_output` in version 1.2 and will be
removed in 1.4. `sparse_output` is ignored unless you leave `sparse`
to its default value.
```

```
warnings.warn(
```

```
c:\Users\Nikshay Jain\AppData\Local\Programs\Python\Python38\lib\site-
packages\sklearn\preprocessing\_encoders.py:808: FutureWarning:
`sparse` was renamed to `sparse_output` in version 1.2 and will be
removed in 1.4. `sparse_output` is ignored unless you leave `sparse`
to its default value.
```

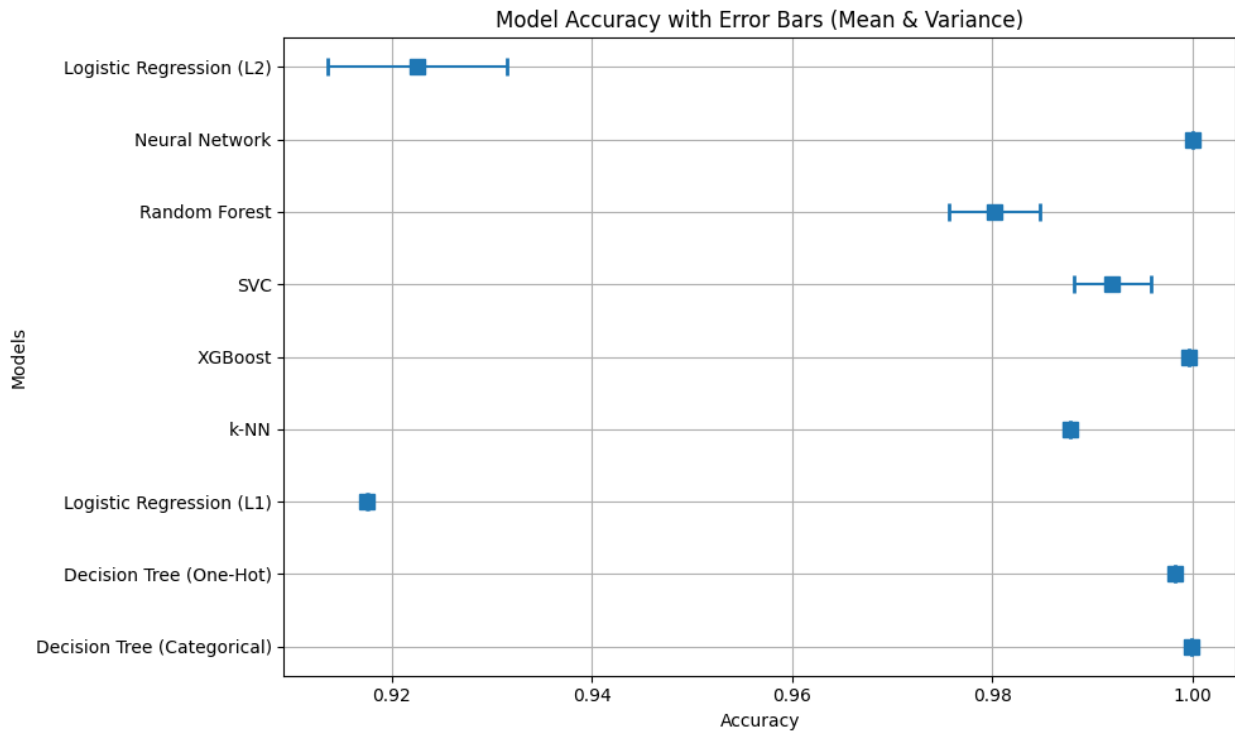
```
warnings.warn(
```

```
c:\Users\Nikshay Jain\AppData\Local\Programs\Python\Python38\lib\site-
packages\sklearn\preprocessing\_encoders.py:808: FutureWarning:
`sparse` was renamed to `sparse_output` in version 1.2 and will be
removed in 1.4. `sparse_output` is ignored unless you leave `sparse`
to its default value.
```

```
warnings.warn(
```

```
C:\Users\Nikshay Jain\AppData\Local\Temp\
ipykernel_3956\1160318684.py:51: UserWarning: marker is redundantly
defined by the 'marker' keyword argument and the fmt string "o" (->
marker='o'). The keyword argument will take precedence.
```

```
plt.errorbar(means, methods, xerr=variances, fmt='o', capsize=5,
capthick=2, marker='s', markersize=8, label='Mean Accuracy')
```



Task 2

```
# Unipolar sigmoid
def unipolar_sigmoid(x):
    return 1 / (1 + np.exp(-x))

# Bipolar sigmoid using unipolar sigmoid
def bipolar_sigmoid(x, a=1):
    return 2 * unipolar_sigmoid(a * x) - 1

def tanh_bipolar_sigmoid(x, a=1):
    return np.tanh(a * x)


a_values = [-5, -1, -0.1, -0.01, 0.001, 0.01, 0.1, 1, 5]
x = np.linspace(-10, 10, 400)

# Plotting the response curves for different 'a' values
for a in a_values:
    y_bipolar = bipolar_sigmoid(x, a)
    y_tanh = tanh_bipolar_sigmoid(x, a)


    plt.plot(x, y_bipolar, label=f'bipolar_sigmoid({a}x)',
             linestyle='--')
    plt.plot(x, y_tanh, label=f'tanh({a}x)', linestyle='--')
    plt.xlabel('x')
    plt.ylabel('y')
```



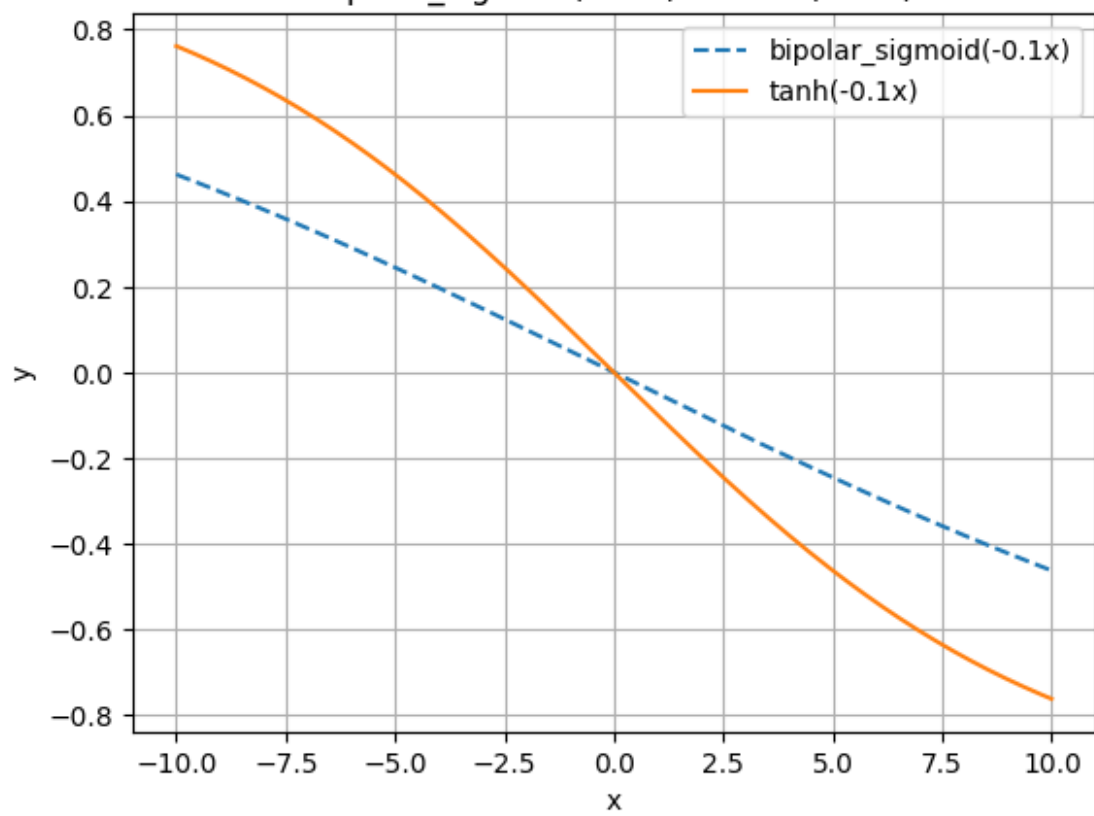
```
plt.title(f'Bipolar_sigmoid({a}x) vs tanh({a}x)')  
plt.legend()  
plt.grid(True)  
plt.show()
```



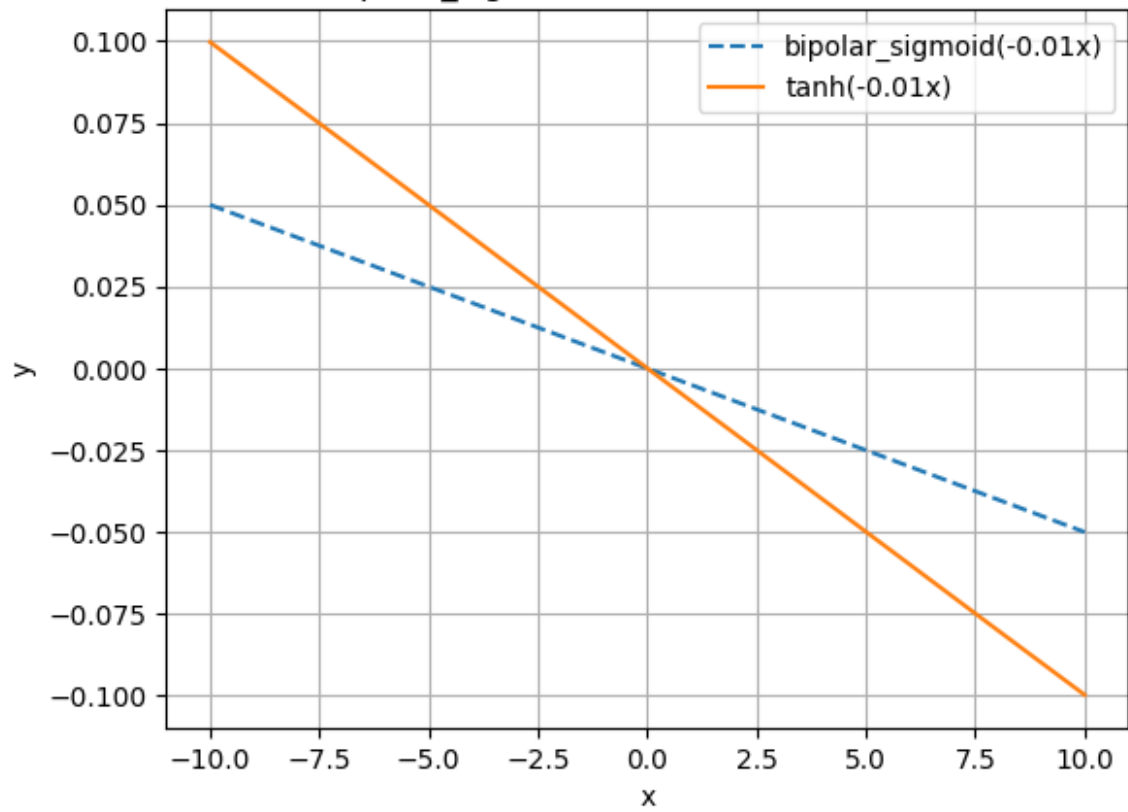
Bipolar_sigmoid(-1x) vs tanh(-1x)

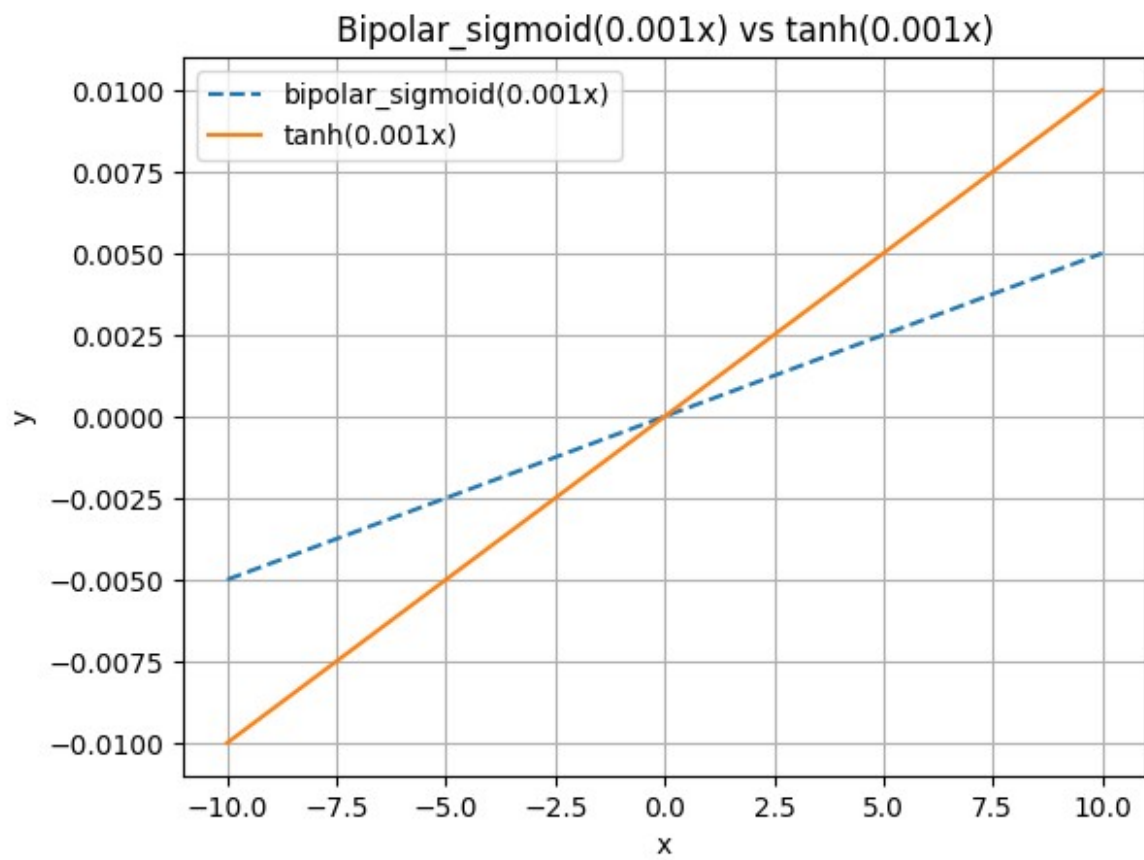


Bipolar_sigmoid(-0.1x) vs tanh(-0.1x)




Bipolar_sigmoid(-0.01x) vs tanh(-0.01x)






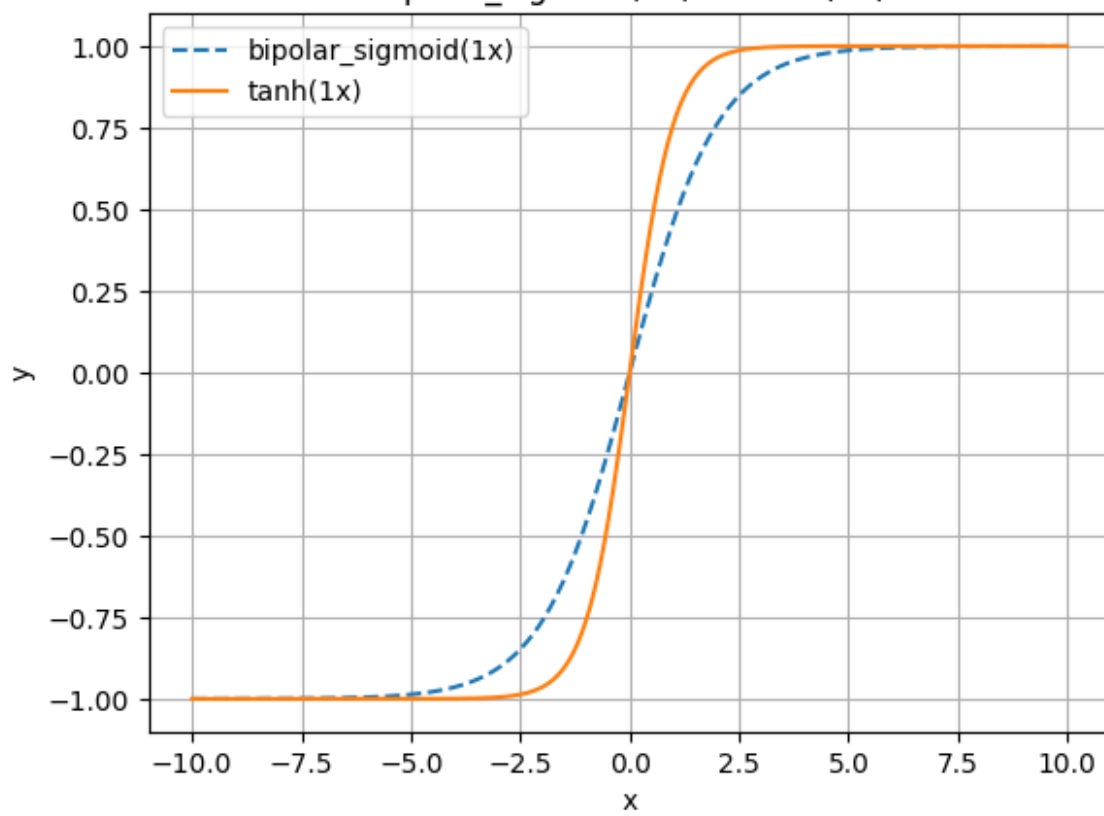
Bipolar_sigmoid(0.01x) vs tanh(0.01x)

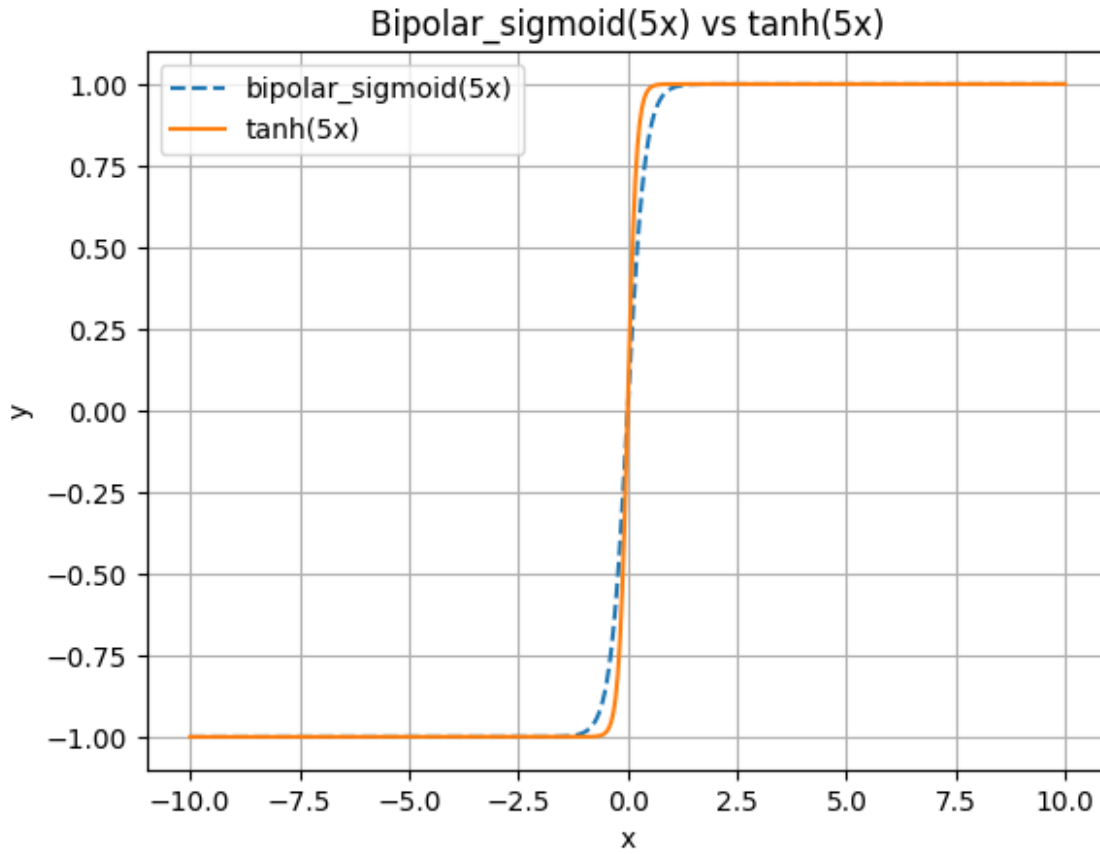


Bipolar_sigmoid(0.1x) vs tanh(0.1x)



Bipolar_sigmoid(1x) vs tanh(1x)





```
# Derivative of bipolar sigmoid
def bipolar_sigmoid_derivative(x):
    return ((1 + bipolar_sigmoid(x))**2) / 2

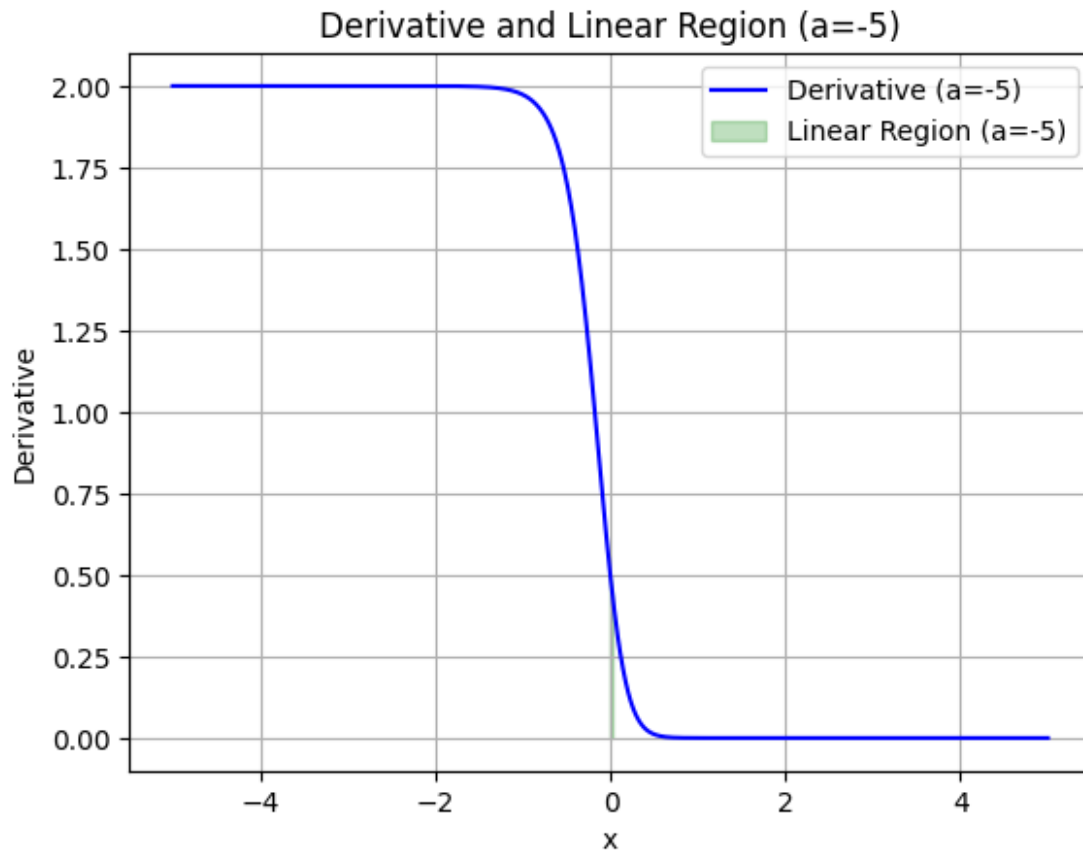
thresh=0.05
x = np.linspace(-5, 5, 1000)

for a in a_values:
    # Compute the derivative of bipolar sigmoid(ax)
    derivative = bipolar_sigmoid_derivative(a*x)

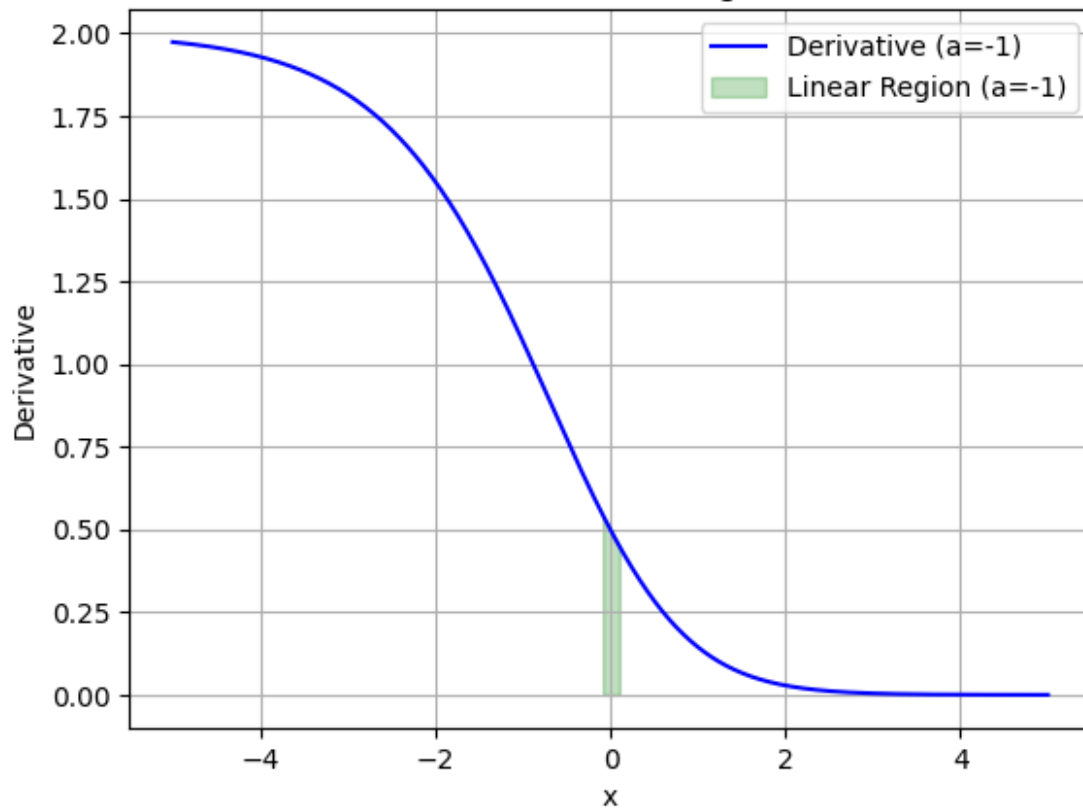
    # Calculate where the derivative is approximately constant (linear
    region)
    center_derivative = derivative[len(derivative) // 2]
    linear_region = np.abs(derivative - center_derivative) < thresh


    # Create a new figure for each value of 'a' and plot the bipolar
    sigmoid derivative and linear region
    plt.plot(x, derivative, label=f'Derivative (a={a})', color='blue')
    plt.fill_between(x, derivative, where=linear_region, color='blue',
    alpha=0.25, label=f'Linear Region (a={a})')
    plt.title(f'Derivative and Linear Region (a={a})')
    plt.xlabel('x')
```


```
plt.ylabel('Derivative')  
plt.legend()  
plt.grid(True)  
plt.show()
```

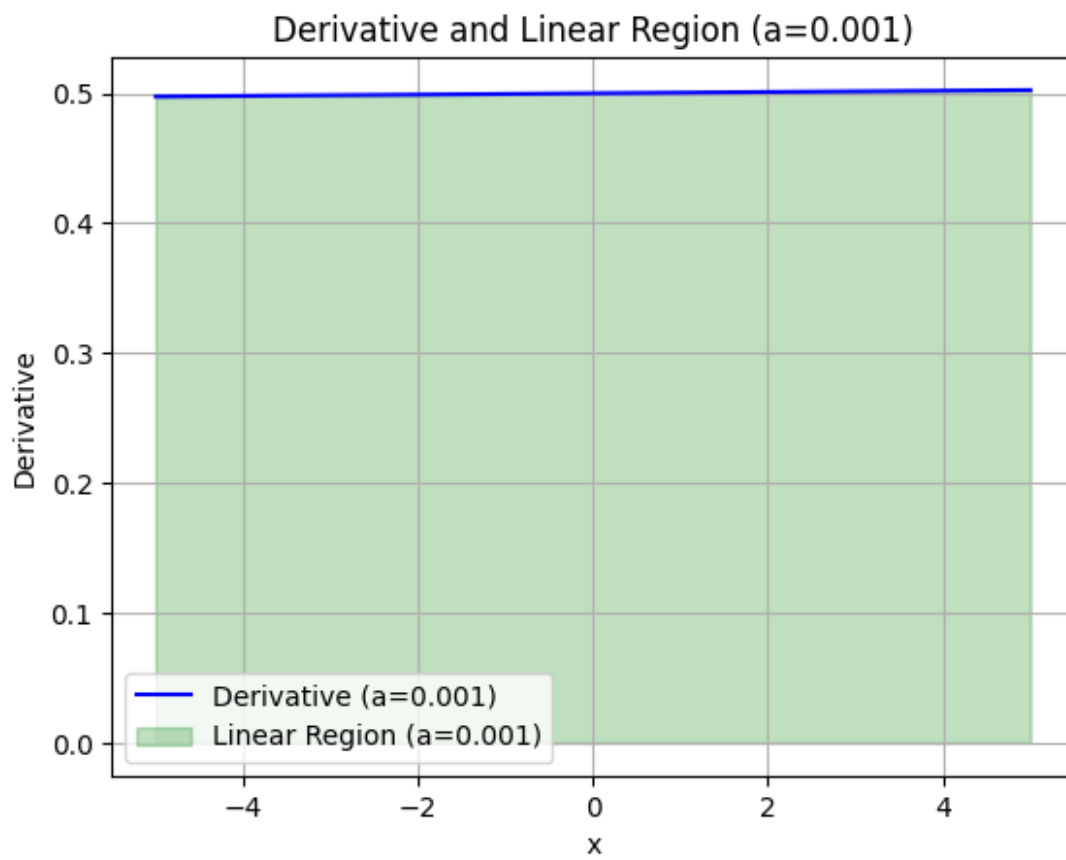


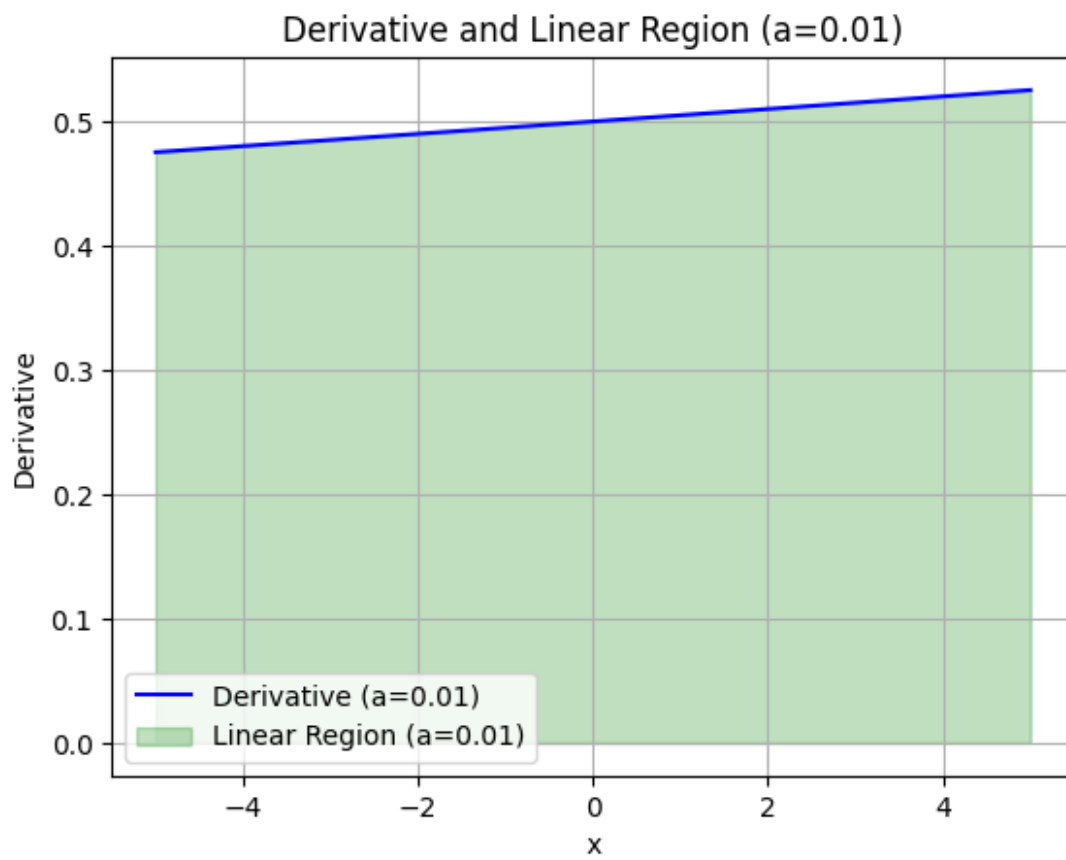
Derivative and Linear Region (a=-1)














Derivative and Linear Region (a=1)

